




FTF | FREESCALE TECHNOLOGY FORUM
POWERING INNOVATION

Multicore Software Migration and Development – Challenges and Solutions

FTF-ENT-F0151


Shrek Wang
Software R&D



August 2012


Freescale, the Freescale logo, AlllVec, C-5, CodeTEST, CodeWarrior, ColdFire, ColdFire+, C-Ware, the Energy Efficient Solutions logo, Kinetis, mobileGT, PowerQUICC, Processor Expert, QorIQ, Qorivva, StarCore, Symphony and VortiQa are trademarks of Freescale Semiconductor, Inc., Reg. U.S. Pat. & Tm. Off. Airfast, BeeKit, BeeStack, CoreNet, Flexis, MagniV, MXC, Platform in a Package, QorIQ Qonverge, QUICC Engine, Ready Play, SafeAssure, the SafeAssure logo, SMARTMOS, TurboLink, Vybrid and Xtrinsic are trademarks of Freescale Semiconductor, Inc. All other product or service names are the property of their respective owners. © 2012 Freescale Semiconductor, Inc.





Agenda


- Objective
- System characteristics of a Multicore Architecture
- Challenges with Multicore Architecture
 - Application Porting Challenges from Unicore to Multicore Architecture
 - Programming Model
 - Processing Model
- Debugging Support with Multicore Architecture
 - Performance Debugging Support
 - Functionality Debugging support
- Summary



Objective

After this presentation, you should be able to understand following:

- Challenges in effectively exploiting the Multicore SoC architecture
- Performance and functionality debugging challenges
 - Effective threading
 - Avoidance of deadlock situations
- Role of OS
 - Thread and interrupt migration across different CPUs
 - Support for thread and interrupt affinity and debugging
- Role of enhanced I/O devices in Multicore migration
 - Load Spreading
 - Data Stashing
- Debugging Support
 - Freescale proprietary debugging support
 - OpenSource debugging support

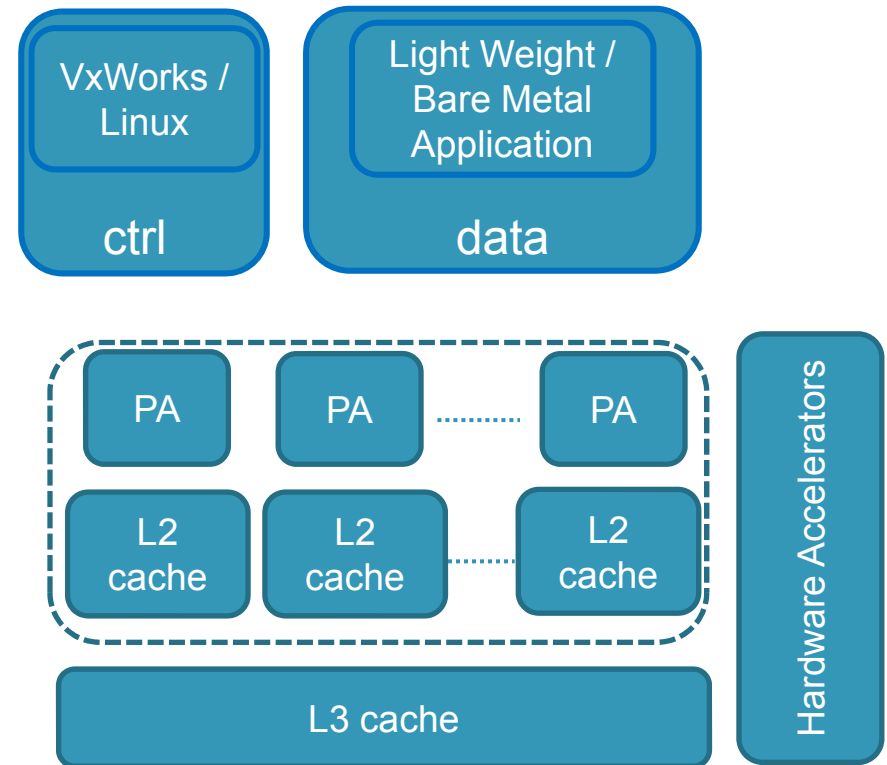


Agenda

- Objective
- System characteristic of a Multicore Architecture
- Challenges with Multicore Architecture
 - Application Porting Challenges from Unicore to Multicore Architecture
 - Programming Model
 - Processing Model
- Debugging Support with Multicore migration
 - Performance Debugging Support
 - Functionality Debugging support

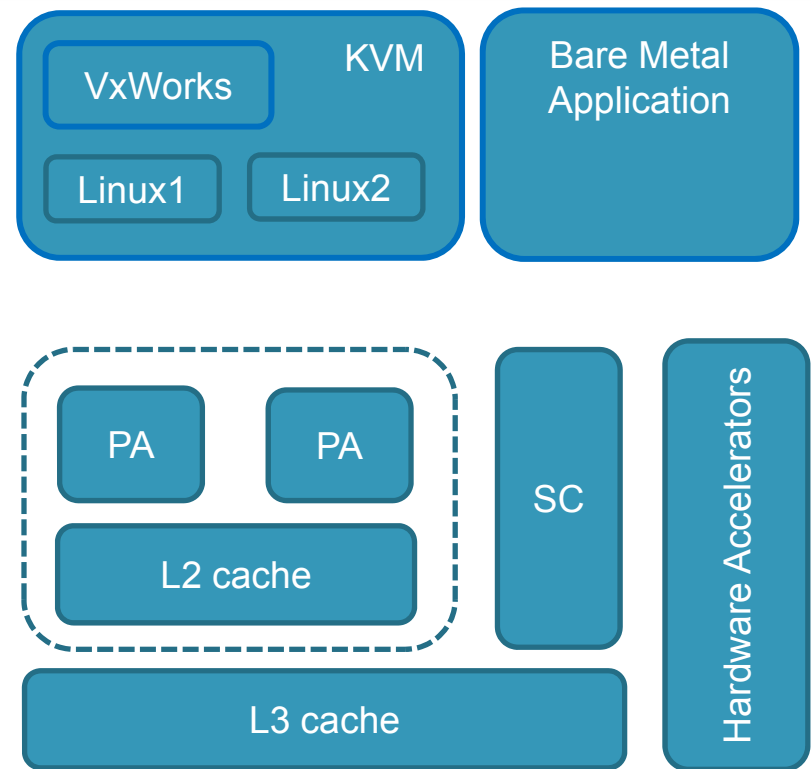
System characteristic changes in Multicore Architecture

- Homogeneous CPUs with multilevel cache hierarchy
 - Threaded cores, co-located control/data planes
 - OSES, baremetal apps running on virtual partitions
- Specialized hardware accelerators, I/O ports with support for load spreading and selective data stashing
- More powerful and complex Buses inside SoC package reduce debug application thread interaction with I/O devices




System characteristic changes in Multicore Architecture (Contd..)

- Heterogeneous CPUs with multilevel cache hierarchy
 - Threaded cores, co-located control/data planes
 - OSes, baremetal apps running on virtual partitions
- Optionally virtual machine appliance (e.g. KVM, etc.)
- More powerful and complex Buses inside SoC package reduce debug application thread interaction with I/O devices



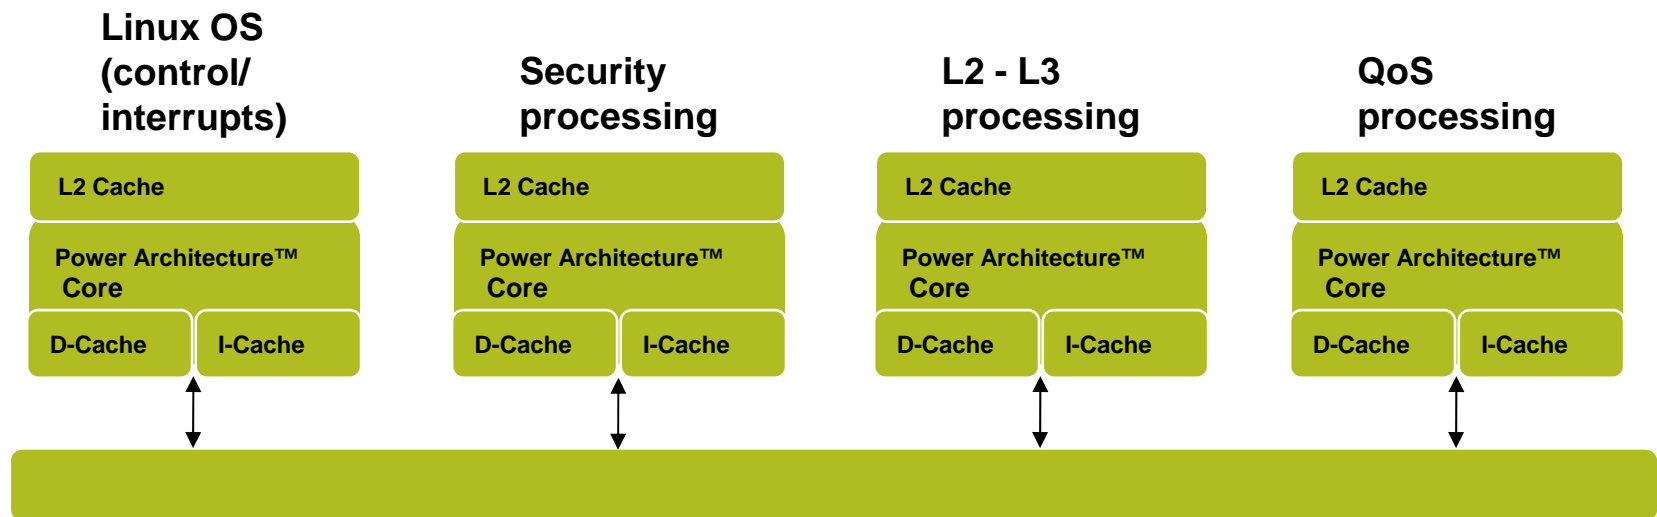
Symmetric Multi-processing (SMP)

- **SMP** is a *multi-processor homogeneous computer architecture* where two or more identical processors are connected to a globally shared main memory...
- Processors could be separate devices, all on 1 device or a mix
- Typically all CPUs share memory, I/O and are run one OS instance
- Each processor can run independent processes and threads
- Any idle processor can be assigned any task

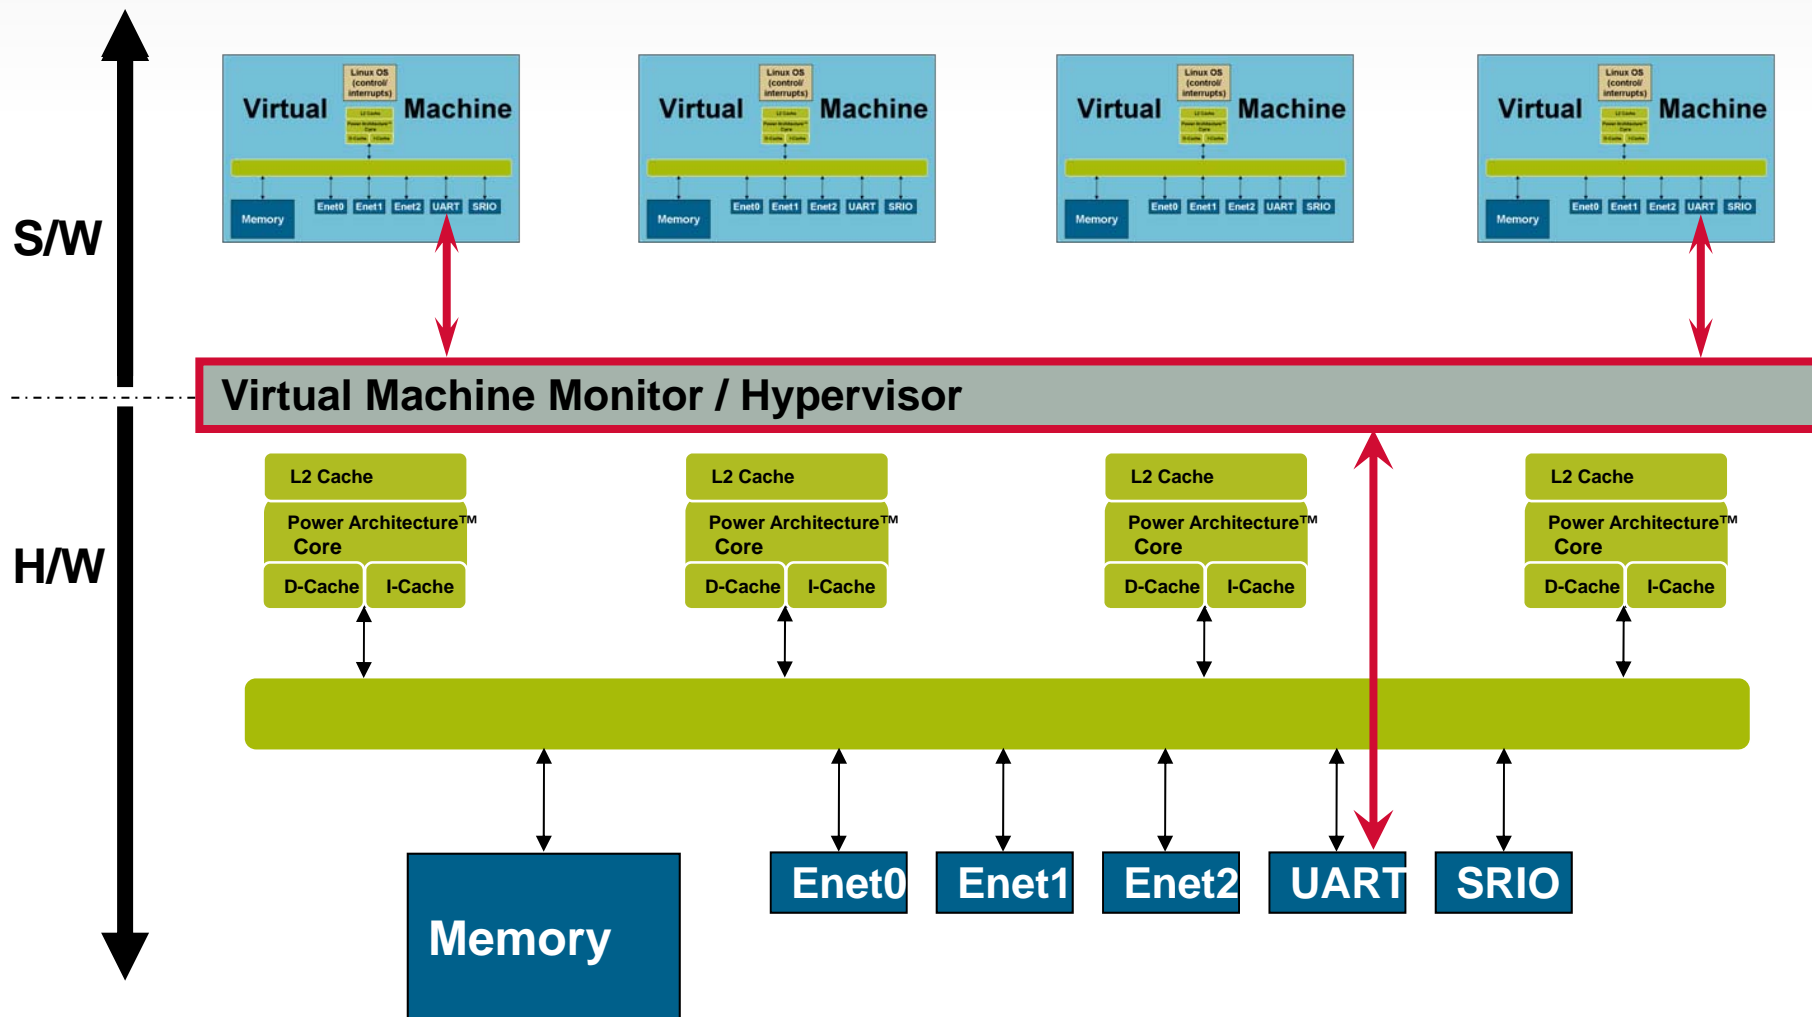


Asymmetric Multi-processing (AMP)

- A *multi-processing usage model* in which individual processors are dedicated to particular tasks, such as running the operating system or performing user requests
- Beware: AMP does not need to imply heterogeneous hardware! It is perfectly possible to implement a software AMP system design on SMP hardware (e.g. MPC8641HPCN Linux BSP supports both)





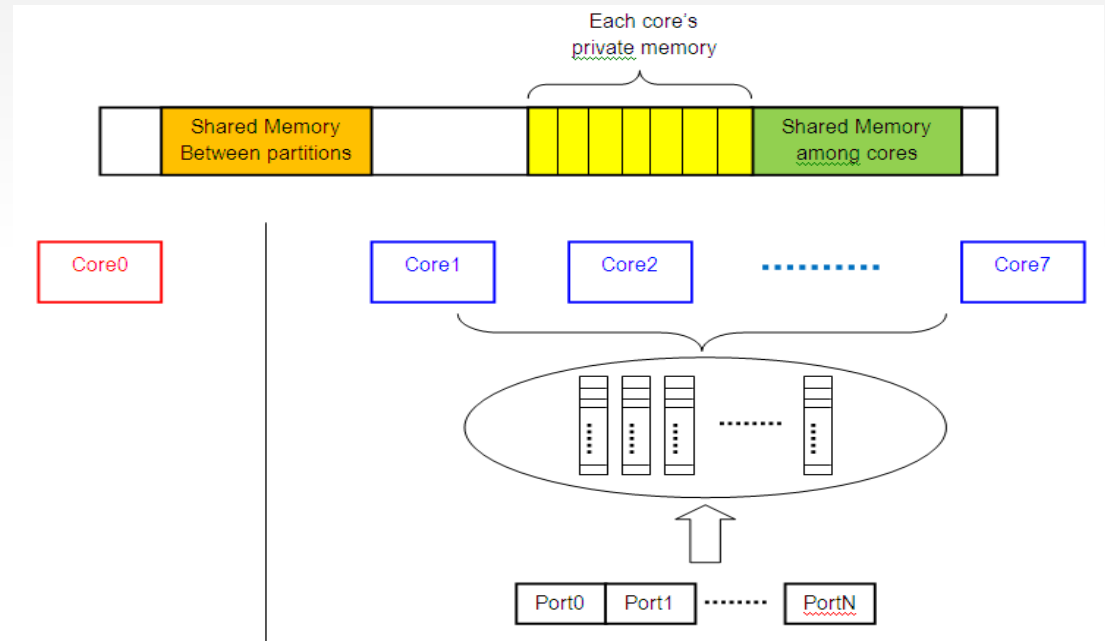
The Need for Virtualization




Software Considerations on Multicore

- System partition (mainly cores, memory, ports resources)
- OS consideration (control plane OS, data plane bare-board or light-weighted env)
- Data plane cores working architecture (functionalities bound to each core/core-group)
- Mutex mechanism
- Data plane tables shared among all data plane cores (shared memory mechanism)
- Inter core communication mechanism
- System global variables, CPU global variables
- Rx/Tx driver
- Arch specific accelerators use
- Control plane partition and data plane partition communications

Examples





Agenda

- Objective
- System characteristic of a Multicore Architecture
- Challenges with Multicore Architecture
 - Application Porting Challenges from Unicore to Multicore Architecture
 - Programming Model
 - Processing Model
- Debugging Support with Multicore migration
 - Performance Debugging Support
 - Functionality Debugging support

Challenges with Multicore Architecture

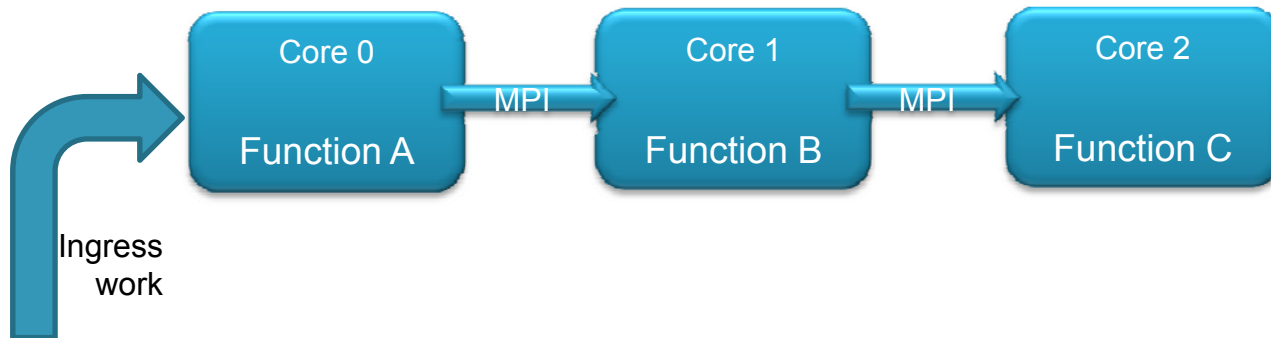
- Legacy software written for non-SMP
 - Multithreading a must for scaling with Multicore
 - Dynamic thread migration by OS scheduler
 - Possibility to Affine application threads and interrupts to a core
 - Locks must for shared resources in multithreaded applications
 - Resource contention, priority inversion, deadlocks
- Parallel threads execution in real time
 - Latent race conditions get uncovered
 - Difficult to reproduce bugs
 - Data corruption more difficult to nail down

Challenges with Multicore Architecture (contd...)

- Cache efficiency requires careful Software design
 - False cache line sharing
 - Cache line thrashing due to threads running on different cores
- Finding performance bottleneck
 - Core loading by scheduler
 - TLB, cache thrashing
 - Stash downgrades
 - Too many interrupts, interrupt distribution, context switches
 - Slow IPC

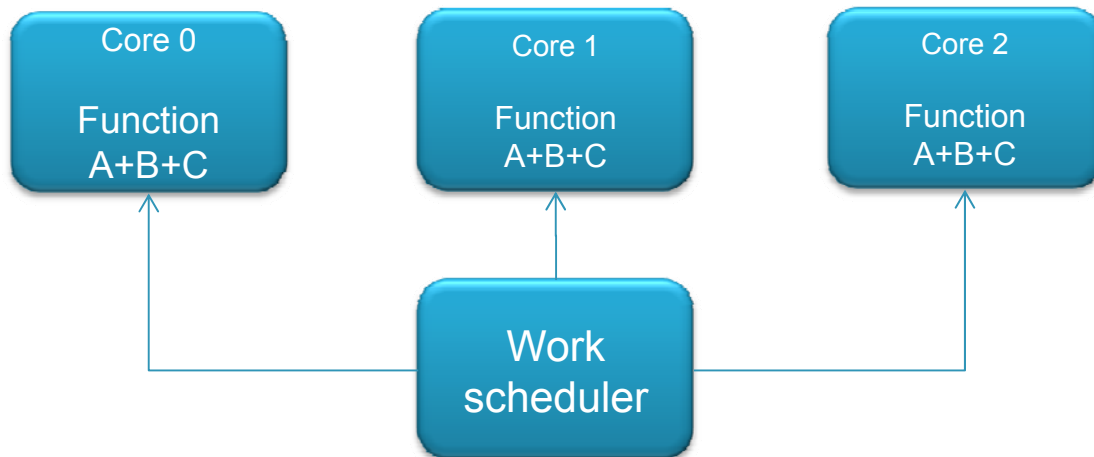
Programming Models

- Pipeline
 - Work division across the cores by the way of functionality
 - Each core executes a subset of functions on each work unit
 - Individual work units passed along functional blocks/cores



Programming Models – contd..

- Run to completion
 - Work is divided by splitting work units across cores
 - Work scheduler divides work units across cores
 - Each core executes all functions on received work units



Programming Models - Comparison

Aspect	Pipeline	Run to completion
Programmability	Difficult to break a problem in equal chunks of work	Easy
I-Cache utilization	Good	Poor
D-cache utilization	Poor	Good
Load balancing	Poor	Easy(If hardware supported)
Fault tolerance	Poor	Good
Locks required	No	Yes
Egress work ordering	Maintained	Disturbed
Work distribution hardware	Not required	Required
IPC/MPI overhead	High	Nil
Real-time/priority work latency	Low	High
Scalability with cores	Difficult (requires redesign)	Easy

Programming Language support

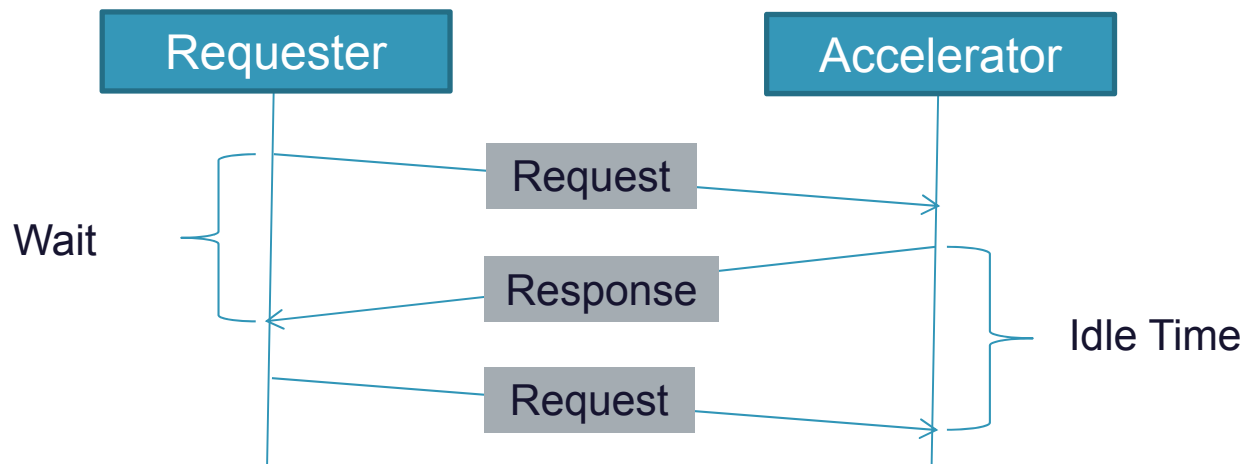
- 'C' does not natively support multithreading
 - Compiler assumes single flow of control
 - Protecting shared resource is responsibility of programmer
 - Locks, volatile etc
 - Parallelism cannot be expressed as an intent to compiler
 - Compiler cannot do inter thread optimization
- Relaxed ordering
 - Out of order completion of loads and stores across cores
 - Instruction/storage barriers
 - Programmers need to be closely aware of hardware

Compiler Limitation

- The compiler may not be able to do the parallelization in the way you like to see it:
 - A loop is not parallelized
 - The data dependency analysis is not able to determine whether it is safe to parallelize or not.
 - The granularity is not high enough
 - The compiler lacks information to parallelize at the highest possible level
- **Explicit parallelization through OpenMP directives and functions are used to overcome above limitations**

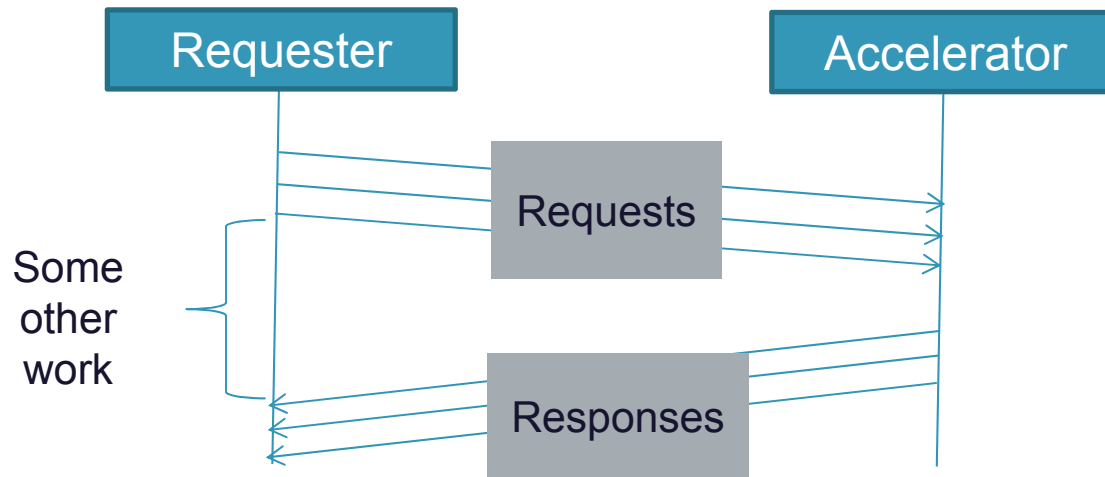
REQ/RESP processing model

- Synchronous
 - Requester waits till response arrives
 - Simplifies programming as response pending state not required
 - Only single pending request
 - Suitable for core-inbuilt offload hardware accelerators




REQ/RESP processing model (contd)

- Asynchronous
 - Requesters do not wait till response arrives (fire and forget)
 - Response comes as another state machine input
 - Request & response co-relation required
 - Multiple requests can be pending
 - Suitable for core external offload hardware accelerators



REQ/RESP processing model - Comparison

Aspect	Synchronous	Asynchronous
Programming ease	High	Low
Context switches	High	Low
Accelerator utilization	Low	High
Pending requests	Single	Multiple
Context switches	High	Low
Single flow B/W	Low	High



Agenda

- Objective
- System characteristic of a Multicore Architecture
- Challenges with Multicore Architecture
 - Application Porting Challenges from Unicore to Multicore Architecture
 - Programming Model
 - Processing Model
- Debugging Support with Multicore migration
 - Performance Debugging Support
 - Functionality Debugging support

OpenSource Debugging Tools

Function trace: Kernel internal tracer

- Derived from -rt patch Latency Tracer
- Plugin tracers
 - ftrace : function tracer
 - irqsoff : interrupt disabled latency
 - wakeup : latency of highest priority task to wake up
 - sched_switch: task context switches
 - (more)
- Works with Ring buffer
- Saved traces used to save maximum latency traces
- **Gives cost-view of different context level in kernel per core**

OpenSource Debugging Tools (Cont.)

Oprofile & Perf

- System-wide profiler (both kernel and user code)
- Sample-based profiler –Time based and event based
- SMP machine support
- Performance monitoring hardware support
- Relatively low overhead, typically <10%
- Designed to run for long times
- Perf supports recording function trace between two events
- **Gives more complete picture of where application spends time per core or process wide**

Freescal Debugging Tools

- **Scenarios Tool**

- Simple one click access to common performance analysis scenarios for QorIQ – including DPAA Debug IP
- Removed need from customer to know how to configure DPAA debug IP to collect performance data from common scenarios.

- **Register Analyzer**

- Alpha version delivered
- Provides feedback on incorrect configuration of DPAA blocks (FM, QM, BM and SEC)

- **SPID (Major Updates for 2012)**

- Target Side Library to access debug IP inside (instrumentation) a customer application.
- Sample code to show how to collect trace and performance data for basic cases.
- Extensive documentation.
- Supports Linux and Bare Metal Applications

- **SPID Utilities (New for 2012)**

- Applications focused on specific use-cases to improve usability of debug IP in FSL SDK Linux applications.
- Packaged for release in SDK 1.2
- First release 3Q 2012

- **Application Notes**


- Performance Analysis
- LTTng Trace
- Application Debugging

- **qorIQ-dbg (Major Updates for 2012)**

- Easy access to QorIQ debug IP from user scripts or applications via standard file system operations.
- Present tools in a manner that is familiar to Linux developers

- **DPAA Tools. (New for 2012)**

- Packet Tracing tool – Allows a customer to trace a packet as it moves through system.
 - Focus for 2012
 - First version scheduled for 4Q 2012.
- This tooling makes use of the on-board trace facility of each block.
- DPAA Trace supporting ALU USDPAA application Alpha being delivered in CW 10.1.2 (may 2012) and Production versions (Expert mode) scheduled for CW 10.2 (June 2012)
- **LTTng/LTTngX**
 - Linux Community standard Kernel Tracing Framework
 - LTTngX – Freescale's unique extensions to LTTng Framework to support access to Freescale QorIQ Debug IP



Summary

- Multicore processor is like a traditional multiprocessor boards shrunk to a single chip. Additionally it may have hardware accelerators
- Application needs to be threaded to benefit from multicore but most legacy applications are single threaded
- To benefit from multicore capabilities, applications need to be careful about many aspects like core / interrupt affinity, cache thrashing, TLB management, priorities, deadlocks, etc.
- There are multiple programming models for multicore applications, with each one having its own advantages and disadvantages
- Several open source tools are available for functionality and performance debugging on multicore platforms
- Freescale supports QorIQ debug architecture to assist debug across layers of software and hardware

Freescale on Kaixin

Tag yourself in photos
and upload your own!



Weibo?

Please use hashtag
#FTF2012#



Session materials will be posted @ www.freescale.com/FTF

